



Delta Lake: Turn your Data Lake into a Lakehouse

Vinoaj (Vinny) Vijeyakumaar
Senior Solutions Architect
vinnv@databricks.com |

©2022 Databricks Inc. — All rights reserved

/vinoaj



Follow along @
<https://vnjv.co/dec22>

lakehouse



★ Rating ▾

⌚ Hours ▾

☰ All filters

Sort by ▾

Lakehouse Cafe

4.9 ★★★★★ (26)

Cafe · 1322 Pittwater Rd

Closed · Opens 5:30AM Mon · 0426

042 683

Dine-in · Kerbside pickup · No-contact delivery



Website



Directions

ORDER TAKEAWAY

The Lakehouse

No reviews

Villa · 27 Woorarra Ave



Directions

The Entrance Lake House

4.3 ★★★★★ (558) · \$\$

Cafe · 27 The Entrance Rd

Breezy locale for contemporary
cuisine

Open · Closes 2PM · (02) 4332

5253

Dine-in · Takeaway · No delivery



Website



Directions

The Lakehouse @ Brightwaters

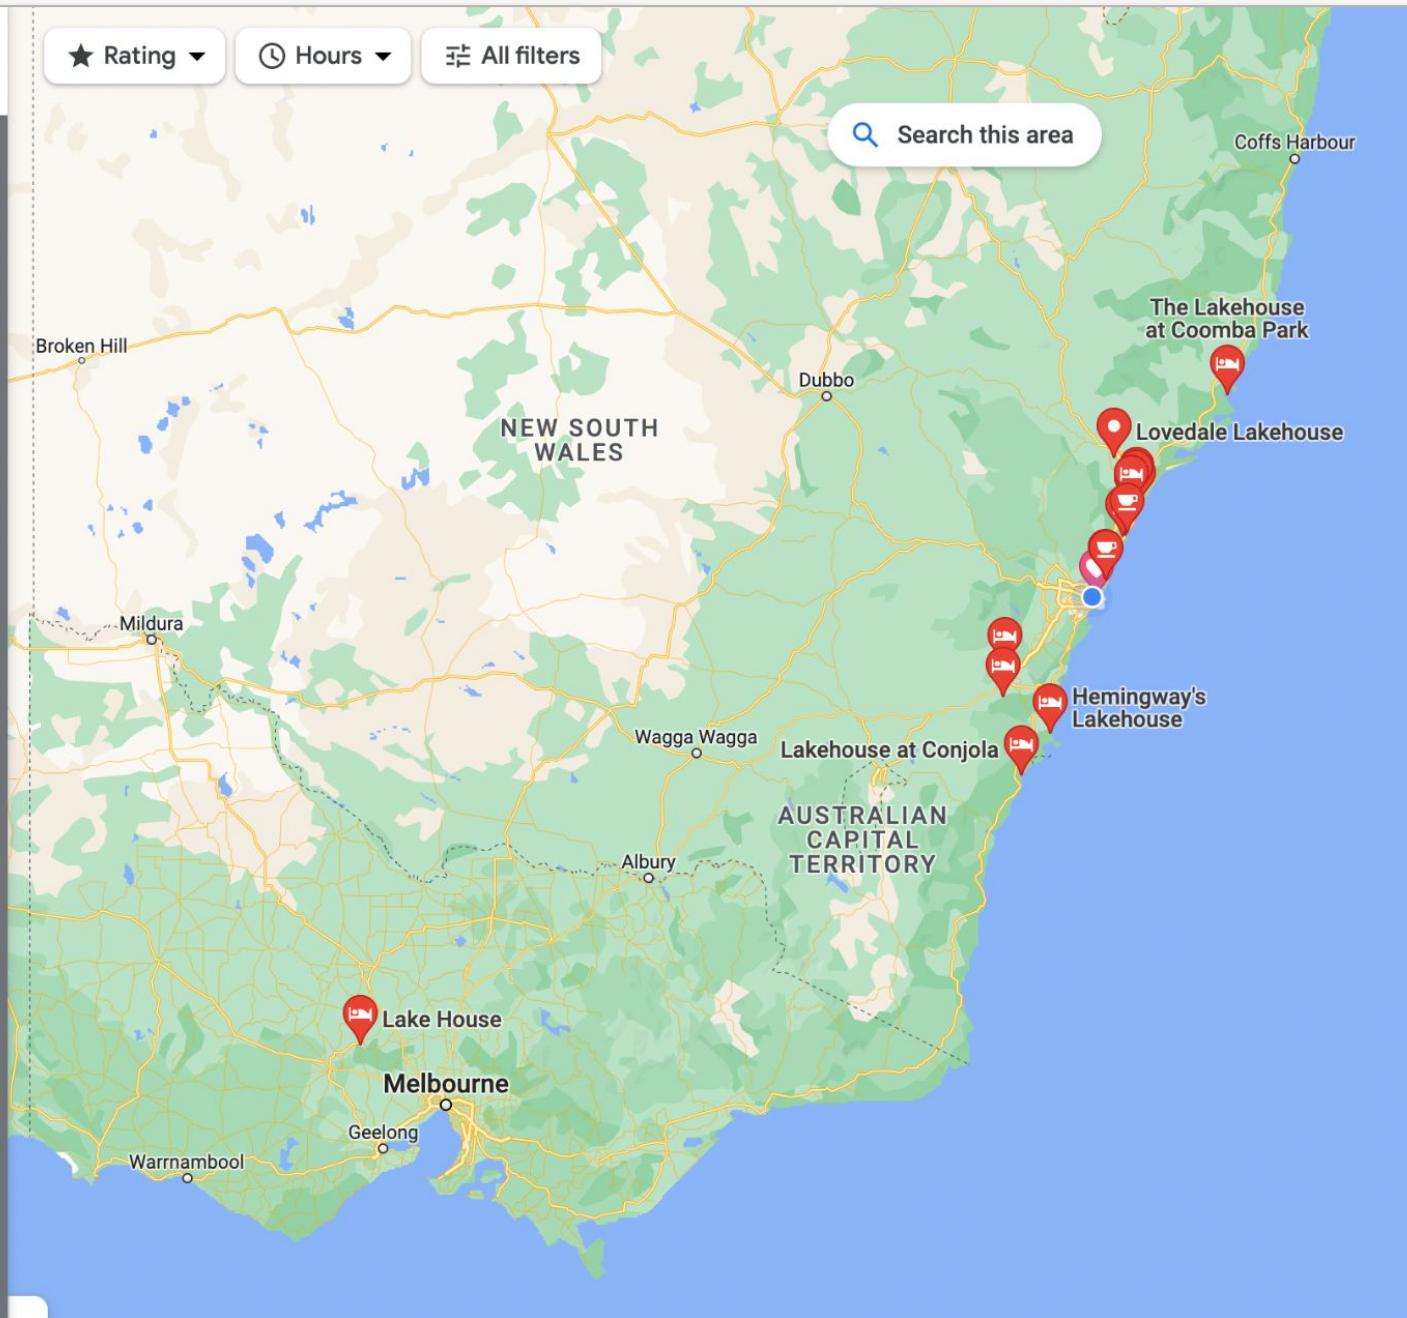
5.0 ★★★★★ (1)

Hotel · 8 Bulgonia Rd

0402 985 598



Directions



> whoami



Vinoaj (**Vinny**) Vijeyakumaar
Senior Solutions Architect

Databricks; Google Cloud; Google Analytics
17+ years in cloud & data space



<https://www.linkedin.com/in/vinoaj/>

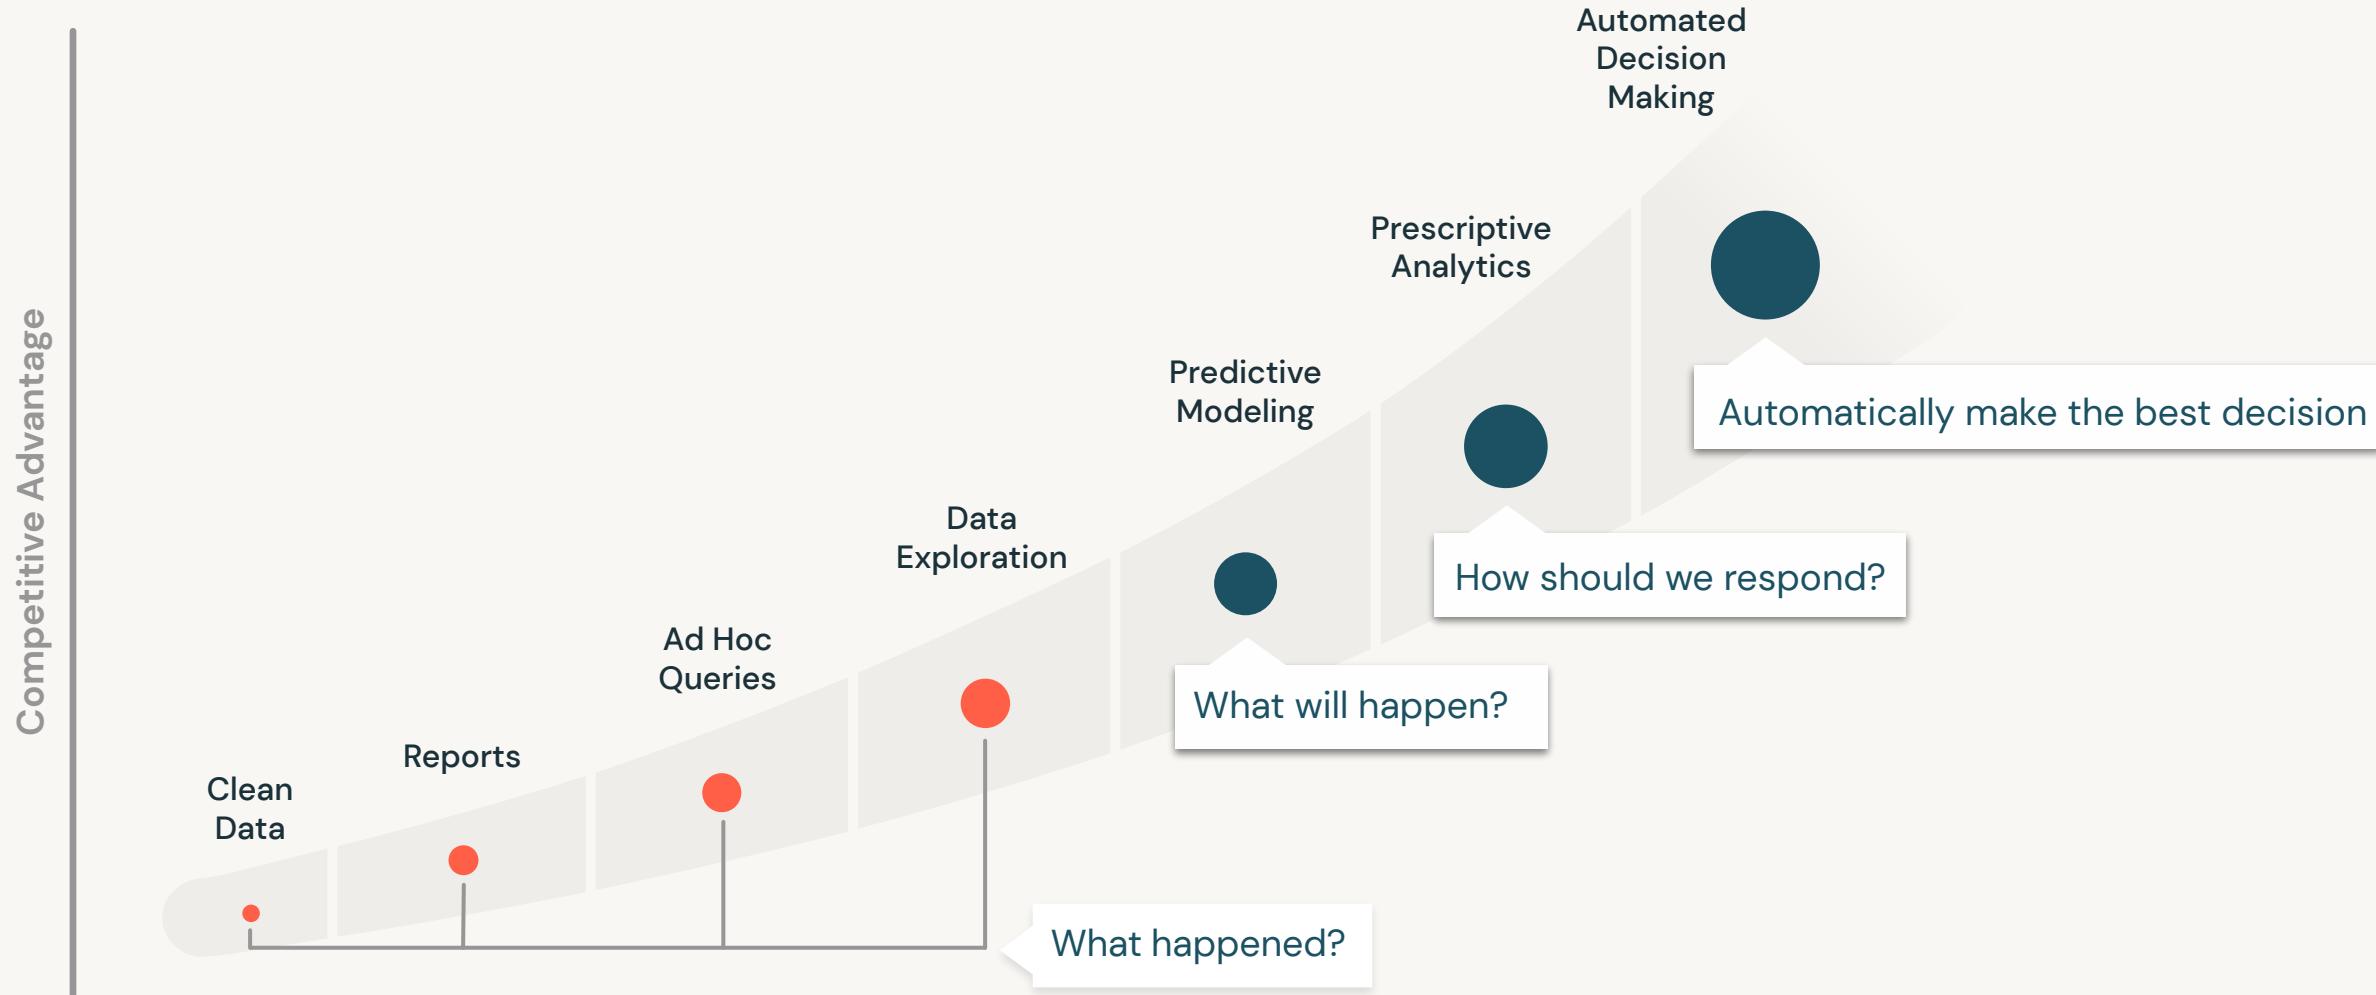
What has driven the need for Lakehouse?

Follow along @
<https://vnjv.co/dec22>

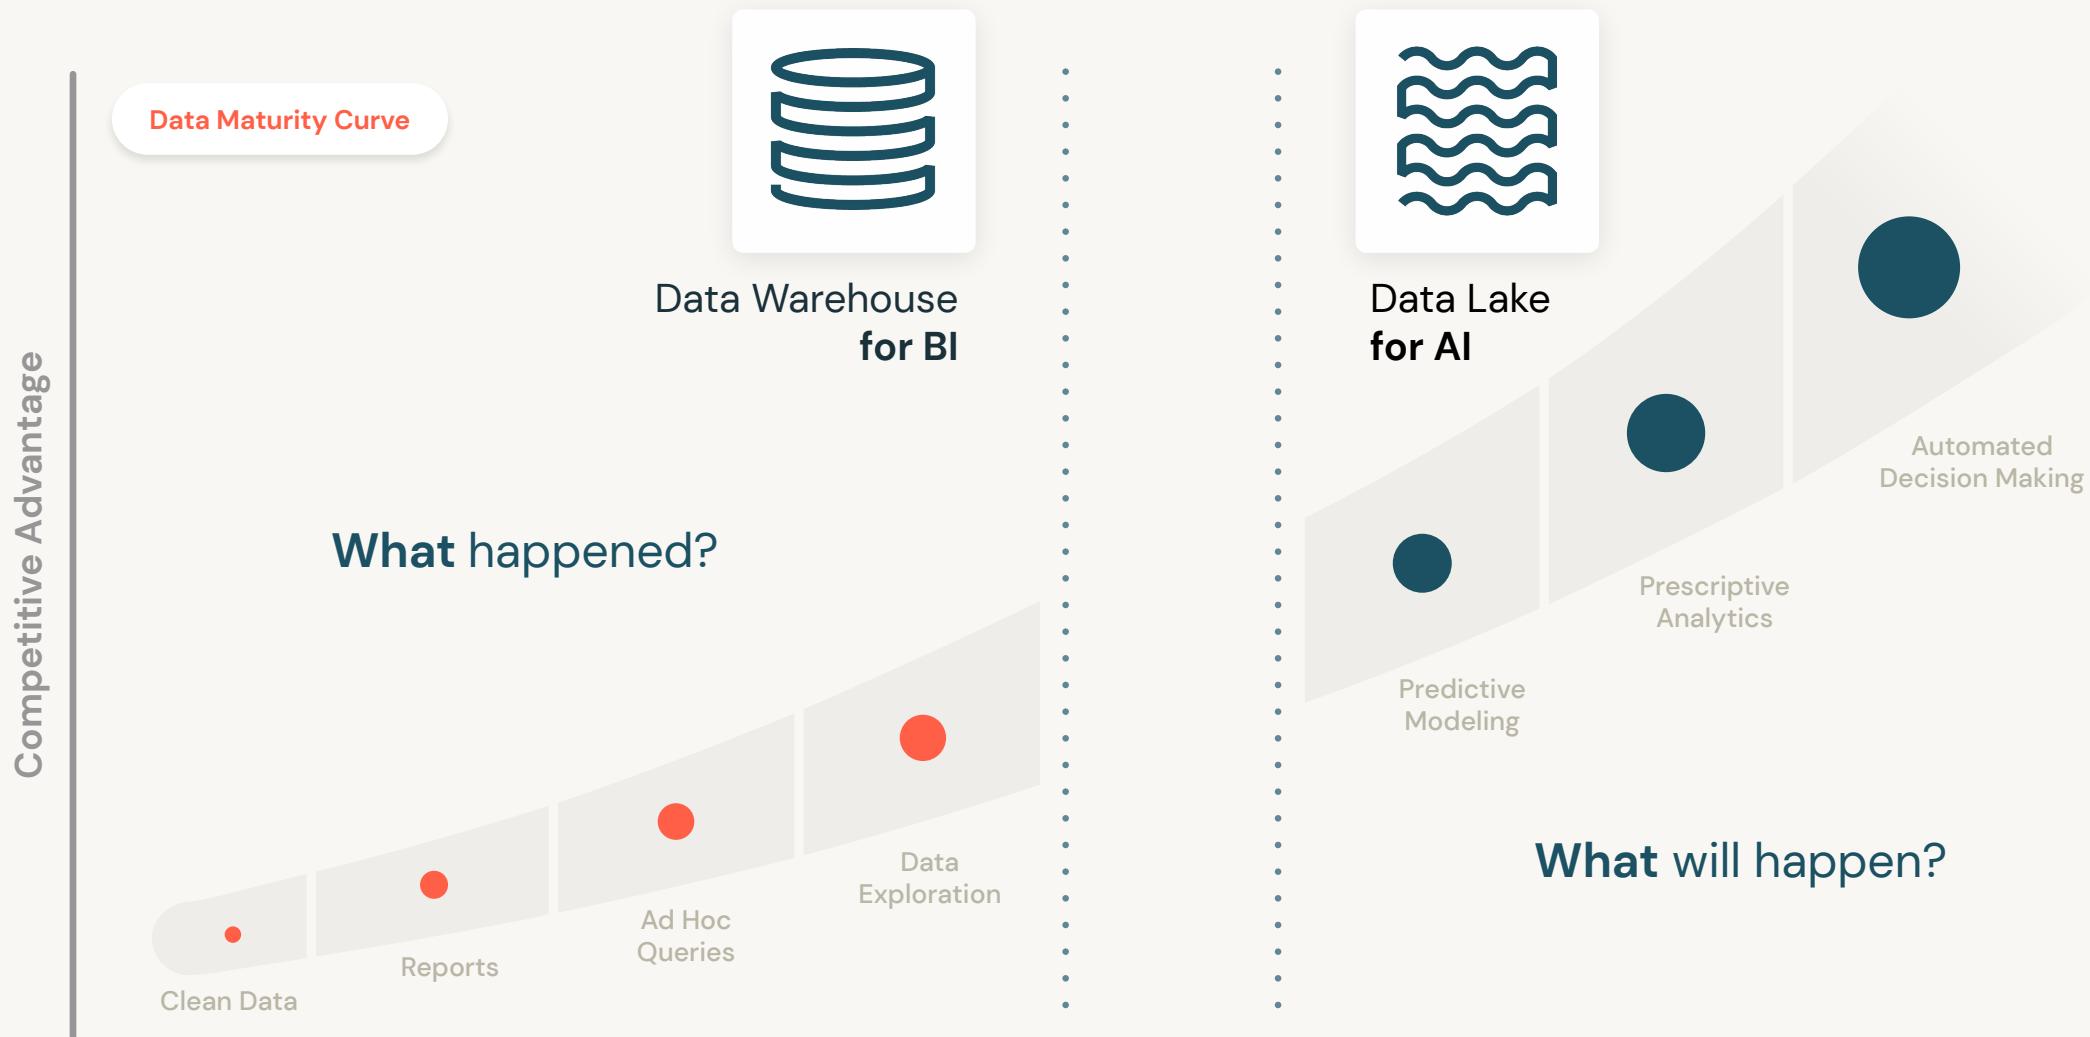


Tech leaders are to the right of the Data Maturity Curve

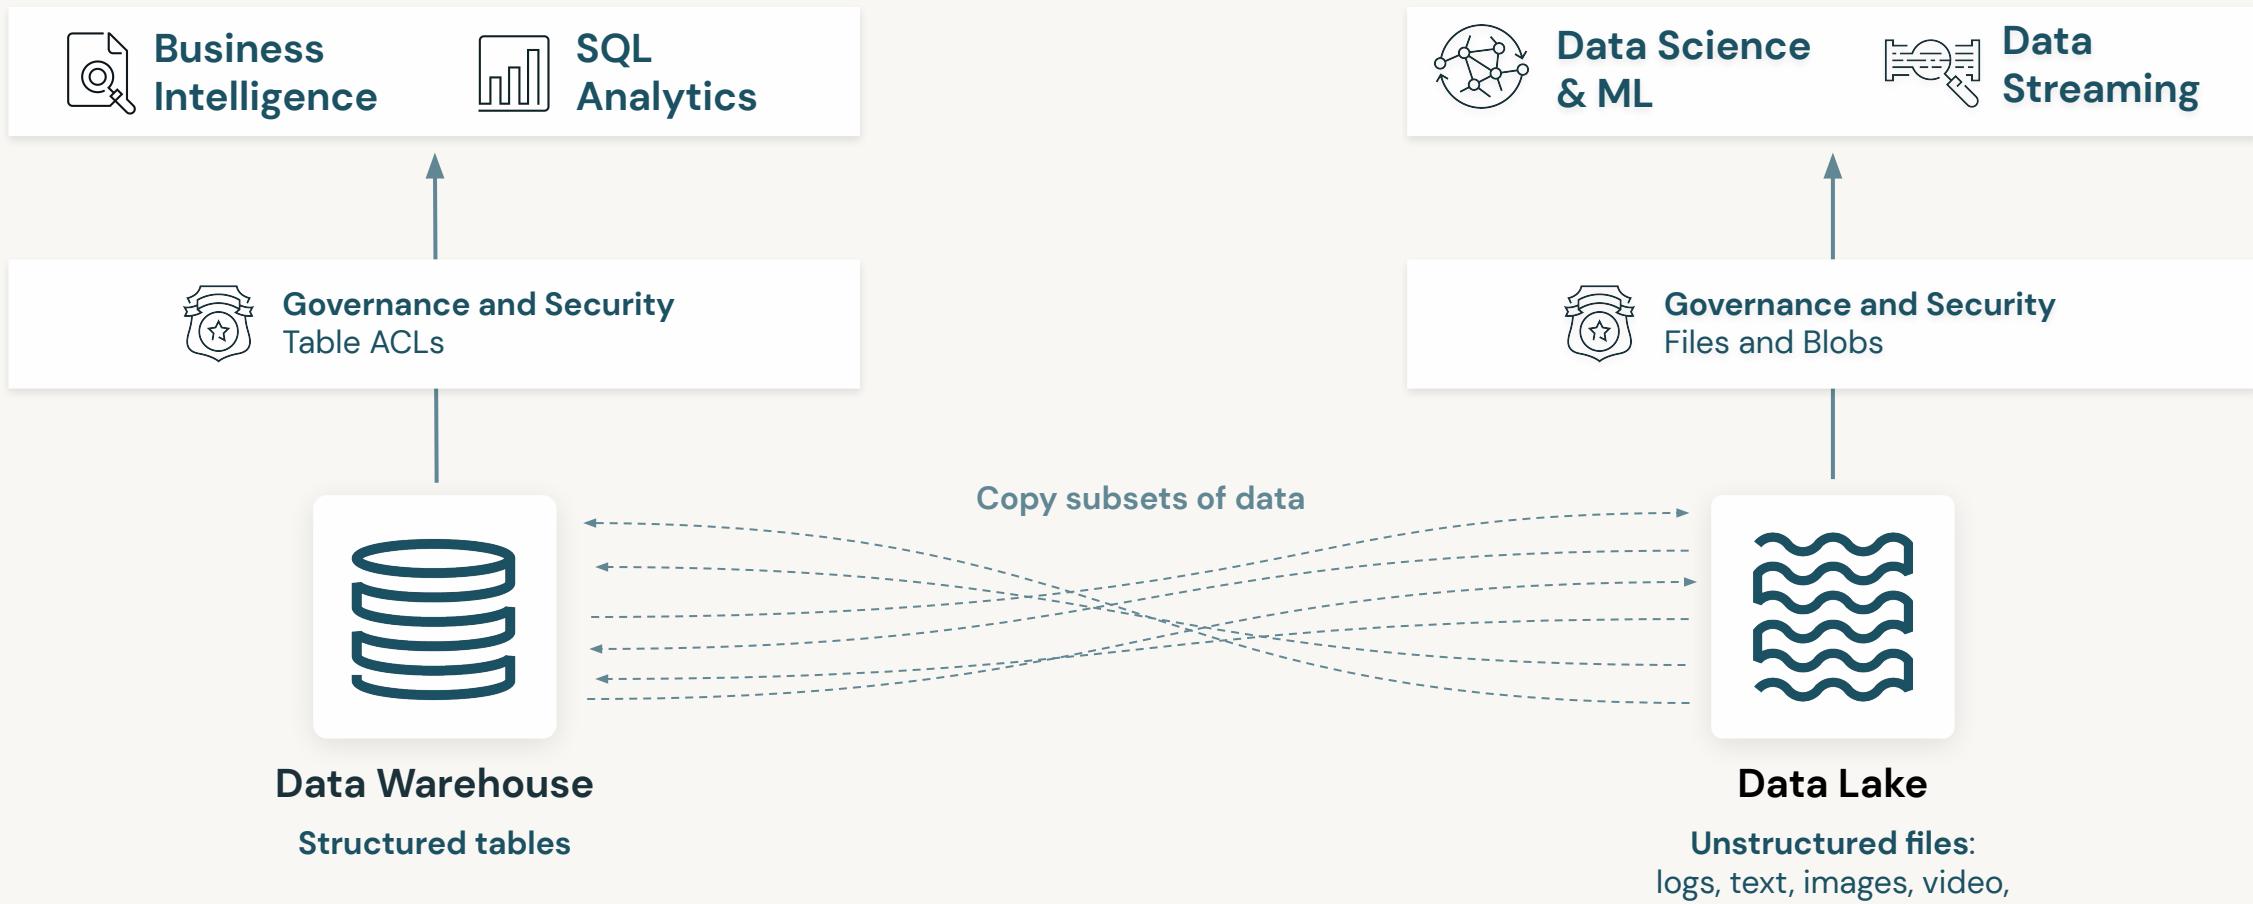
From hindsight to foresight



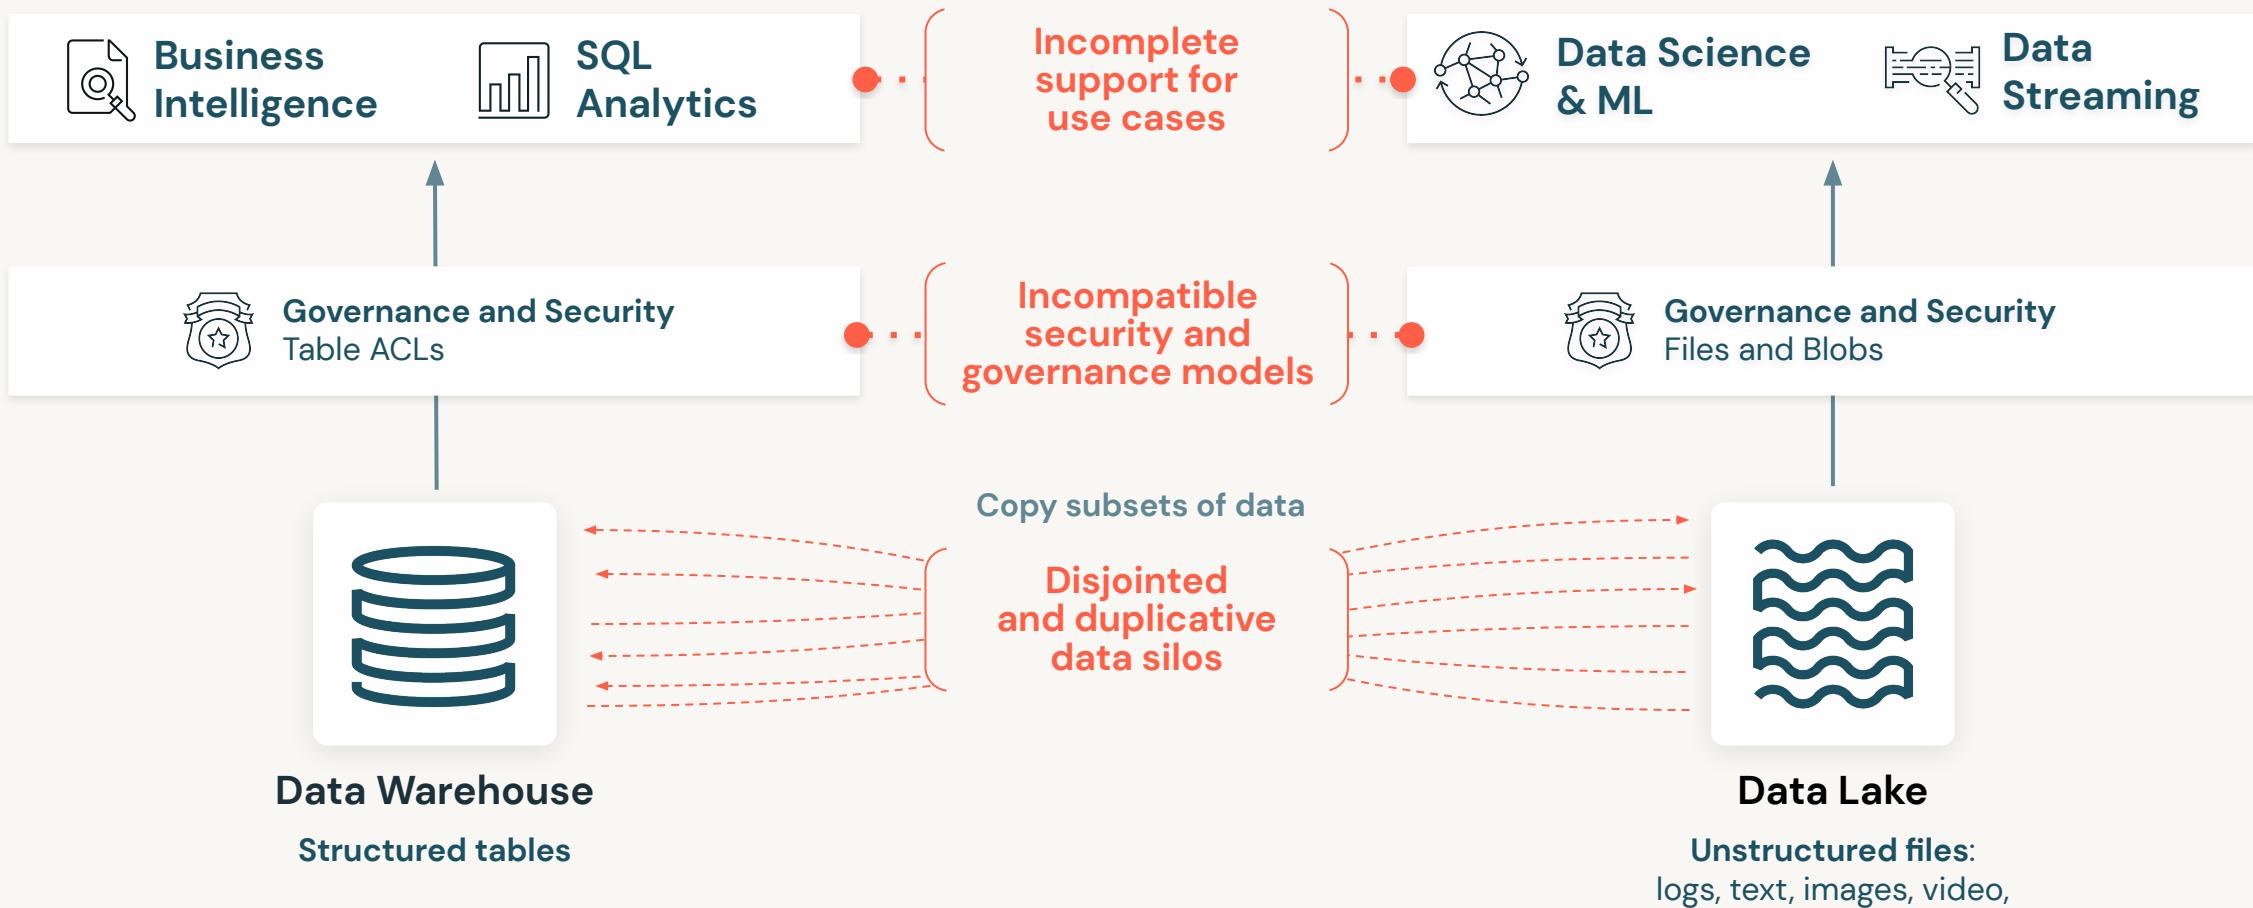
Realizing this requires two disparate, incompatible data platforms



Realizing this requires two disparate, incompatible data platforms



Realizing this requires two disparate, incompatible data platforms



Realizing this requires two

compatible data platforms



Lakehouse Platform

All machine learning, SQL, BI,
and streaming use cases

One security and governance
approach for all data assets on all clouds

An open and reliable data platform
to efficiently handle all data types

Incomplete support for use cases

Incompatible security and governance models

Disjointed and duplicative data silos

Data Warehouse

Structured tables

Governance



Data Streaming

Governance and Security and Blobs



Data Lake

Unstructured files:
logs, text, images, video,



Delta Lake makes the Lakehouse possible

Follow along @
<https://vnjv.co/dec22>



Integrations



presto



trino



kafka



Azure Synapse



DataHub



Coming Soon:



Streaming



Batch



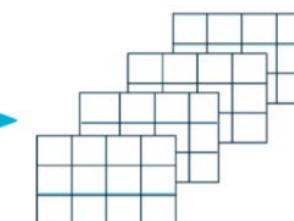
DELTA LAKE



Ingestion Tables
(Bronze)



Refined Tables
(Silver)



Feature/Agg Data Store
(Gold)

Analytics
and Machine
Learning



Azure
Data Lake Storage



amazon
S3



IBM Cloud



THE
LINUX
FOUNDATION

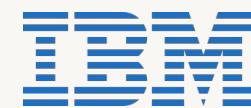
Community-driven!



BackMarket



COMCAST



SAFE GRAPH

SAMBA TV



Starburst



talend*



UCDAVIS



20 Jun 2019



liwensun



v0.2.0

-o ae3daa8

Compare ▾

Delta Lake 0.2.0

We are delighted to announce the availability of Delta Lake 0.2.0!

To try out Delta Lake 0.2.0, please follow the [Delta Lake Quickstart](#).

This release introduces two main features:

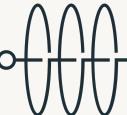
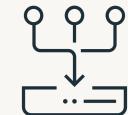
- **Cloud storage support** - In addition to HDFS, you can now configure Delta Lake to read and write data on cloud storage services such as Amazon S3 (issue [#39](#)) and Azure Blob Storage (issue [#40](#)). See [here](#) for configuration instructions.
- **Improved concurrency** (issue [#69](#)) - Delta Lake now allows concurrent *append-only* writes while still ensuring serializability. To be considered as append-only, a writer must be only adding new data without reading or modifying existing data in any way. See [here](#) for more details.

We have also greatly expanded the test coverage as part of this release.



Delta Lake 2.0 – All of Delta is now open

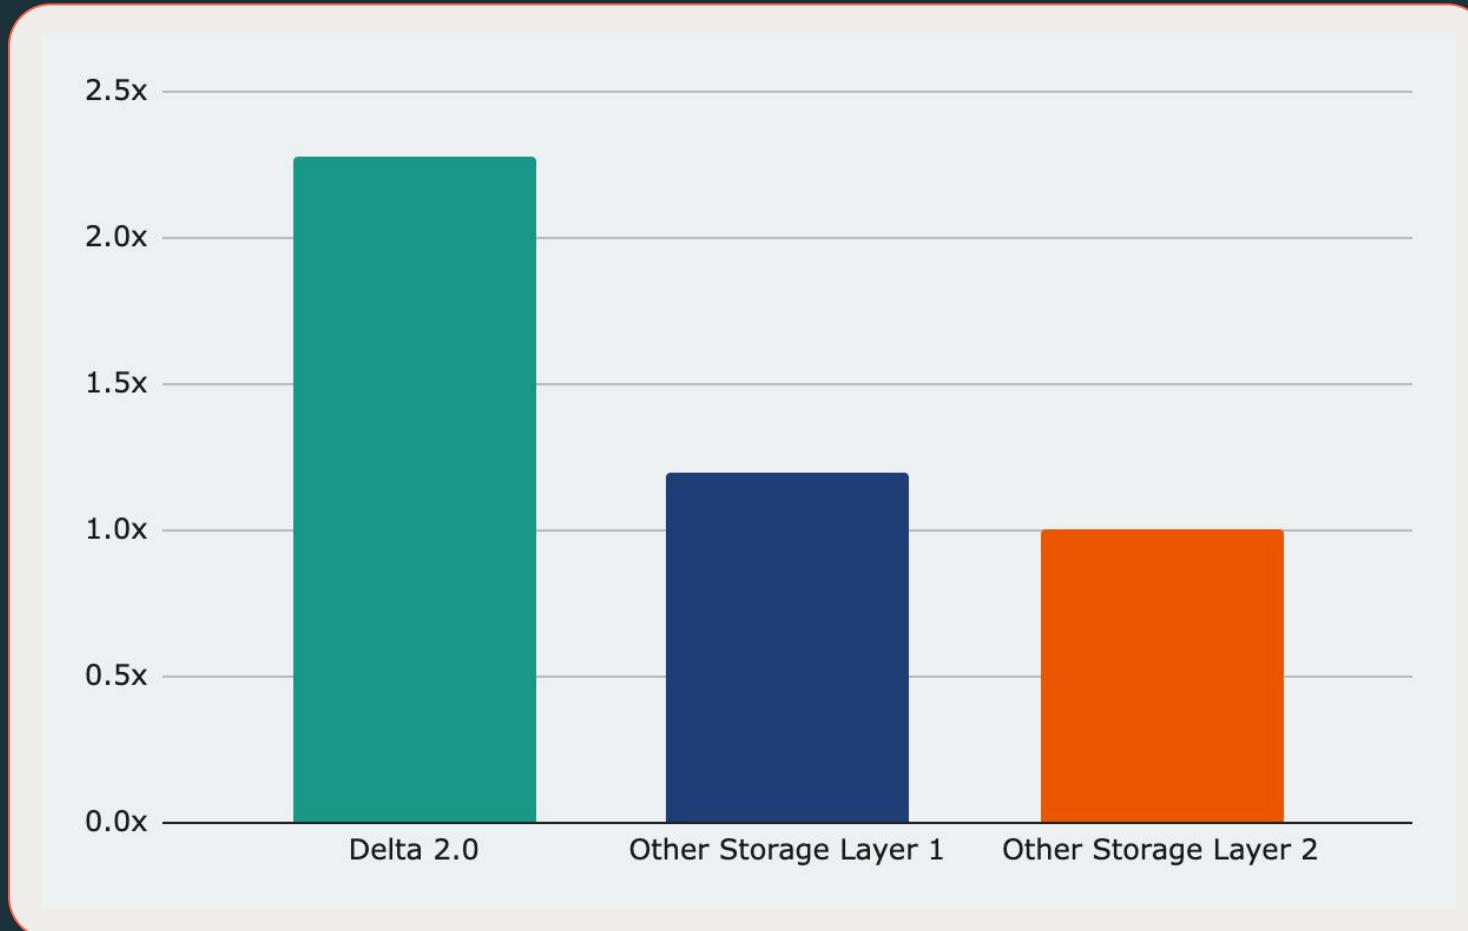
Unlock the power of Delta Lake

								
ACID Transactions	Scalable Metadata	Time Travel	Open Source	OPTIMIZE	OPTIMIZE ZORDER	Change data feed	Generated column support w/ partitioning	Coming Soon!
								
Unified Batch/Streaming	Schema Evolution /Enforcement	Audit History	DML Operations	Table Restore	S3 Multi-cluster writes	Data Skipping via Column Stats	Identity Columns	Iceberg to Delta converter
								
Compaction	MERGE Enhancements	Stream Enhancements	Simplified Logstore	Generated Columns	Multi-part checkpoint writes	Column Mapping	Subqueries in deletes and updates	Fast metadata only deletes



Delta Lake performance

TPC-DS Benchmark Comparison (Higher is better)



Delta Lake is
1.9x faster than
Storage Format 1
2.3x faster than
Storage Format 2



Delta Lake: My Favourite Bits

(that fit into 10 minutes)



Sample Data: Kaggle Stack Overflow Competition

Featured Prediction Competition

Predict Closed Questions on Stack Overflow

Predict which new questions asked on Stack Overflow will be closed

Stack Overflow · 161 teams · 10 years ago

[Predict Closed Questions on Stack Overflow](#)

Follow along @
<https://vnjv.co/dec22>



train.csv (3.73 GB)										
Detail	Compact	Column								10 of 15 columns
↳ PostId	↳ PostCreationDate	↳ OwnerUserId	↳ OwnerCreationDate	# ReputationAtPostCre...	# OwnerUndeletedAns...	Δ Title	Δ BodyMarkdown	3363198 unique values	33702 unique val	
4	11.8m	1Aug08	1Aug12	1	1.57m					
4	07/31/2008 21:42:52	8	07/31/2008 21:33:24	1	0	Decimal vs Double?	I'm new to C# want to use a trackbar for forms opacity is my code de tr...			
6	07/31/2008 22:08:08	9	07/31/2008 21:35:26	1	0	Percentage width child in absolutely positioned parent doesn't work in IE7	I've got an absolutely positioned di containing se children, one which is a relatively po			
8	07/31/2008 23:33:19	9	07/31/2008 21:35:26	16	1	Tools for porting J# code to C#	Are there any conversion to porting Visua code to C#?			
9	07/31/2008 23:40:59	1	07/31/2008 14:22:31	1	1	How do I calculate someone's age in c#?	Given a DateT representing birthday, how calculate som age?			
9610539	03/07/2012 23:07:09	1021610	10/31/2011 08:26:49	29	0	retrieve data from NSUserDefaults to UITableView	I save values labels through NSUserDefault (IBAction) saveDat...			
39	08/01/2008 12:43:11	33	08/01/2008 12:32:10	1	0	Reliable Timer in a Console Application	I am aware in net there are timer types ([http://msdn.ft.com/en-us/magazine/c			

Let's create our first Delta Table

```
● ● ●

filename = "train.csv"
filepath =
    f"/Users/{USERNAME}/custom_demos/input_files/{KAGGLE_COMPETITION}/{filename}"
print(f"Loading contents of: {filepath}")

data_csv = (spark.read
            .option("header", "true")
            .option("multiLine", "true")
            .option("escape", "'') # Handle double quotes in multiline entries
            .csv(filepath))

(data_csv.write
    .format("delta")
    .mode("overwrite"))
# All Delta Lake files (i.e. data and log files) for this table will be stored here
    .save(f"/Users/{USERNAME}/custom_demos/lakehouse/stackoverflow_train00"))
```



```
(spark.read.format("delta")
    .load(f"/Users/{USERNAME}/custom_demos/lakehouse/stackoverflow_train00")
    .display())
```

Table +

	PostId	PostCreationDate	OwnerUserId	OwnerCreationDate	ReputationAtPostCreation	OwnerUndeletedAnswerCountAtPostTime	Title
1	4	07/31/2008 21:42:52	8	07/31/2008 21:33:24	1	0	Decimal vs Double?
2	6	07/31/2008 22:08:08	9	07/31/2008 21:35:26	1	0	Percentage width child in absolutely positioned par
3	8	07/31/2008 23:33:19	9	07/31/2008 21:35:26	16	1	Tools for porting J# code to C#
4	9	07/31/2008 23:40:59	1	07/31/2008 14:22:31	1	1	How do I calculate someone's age in c#?
5	9610539	03/07/2012 23:07:09	1021610	10/31/2011 08:26:49	29	0	retrieve data from NSUserDefaults to TableView
6	39	08/01/2008 12:43:11	33	08/01/2008 12:32:10	1	0	Reliable Timer in a Console Application

Truncated results, showing first 1,000 rows. ▾ | 3.23 seconds runtime

Create a Delta Lake table using SQL

```
CREATE OR REPLACE TEMPORARY VIEW tvw_train
USING CSV -- Databricks-specific operator
OPTIONS(
    path "/Users/vinny.vijeyakumaar@databricks.com/custom_demos/input_files/predict-closed-questions-on-stack-
overflow/train.csv",
    header "true", multiLine "true", escape ''
);

-- Create the Delta table
CREATE OR REPLACE TABLE vinny_vijeyakumaar.stackoverflow_train01
USING DELTA
PARTITIONED BY (PostCreationDate)
TBLPROPERTIES(
    delta.enableChangeDataFeed = true,
    -- autoOptimize.* available in Databricks; P1 priority on current Delta Lake OSS roadmap
    delta.autoOptimize.optimizeWrite = true,
    delta.autoOptimize.autoCompact = true
)
-- (Optional) All Delta Lake files (i.e. data and log files) for this table will be stored here
-- If not specified, location will be managed by your active metastore
LOCATION "/Users/vinny.vijeyakumaar@databricks.com/custom_demos/lakehouse/stackoverflow_train01"
AS
SELECT
    PostId
    , DATE(TO_TIMESTAMP(PostCreationDate, "MM/dd/yyyy HH:mm:ss")) AS PostCreationDate
    , * EXCEPT (PostId, PostCreationDate)
FROM tvw_train;

-- Cleanup
DROP VIEW IF EXISTS tvw_train;

-- Show us the data
SELECT *
FROM vinny_vijeyakumaar.stackoverflow_train01;
```

Follow along @
<https://vnjv.co/dec22>



Let's take a closer look

```
CREATE OR REPLACE TABLE vinny_vijeyakumaar.stackoverflow_train01
USING DELTA
PARTITIONED BY (PostCreationDate)
TBLPROPERTIES(
    delta.enableChangeDataFeed = true,
    -- autoOptimize.* available in Databricks; P1 priority on current Delta Lake OSS roadmap
    delta.autoOptimize.optimizeWrite = true,
    delta.autoOptimize.autoCompact = true
)
-- (Optional) All Delta Lake files (i.e. data and log files) for this table will be stored here
-- If not specified, location will be managed by your active metastore
LOCATION "/Users/vinny.vijeyakumaar@databricks.com/custom_demos/lakehouse/stackoverflow_train01"
AS
SELECT
```





```
SELECT *
FROM vinny_vijeyakumaar.stackoverflow_train01;
```

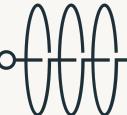
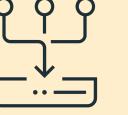
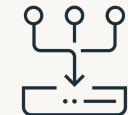
Table + Give feedback

	PostId	PostCreationDate	OwnerUserId	OwnerCreationDate	ReputationAtPostCreation	OwnerUndeletedAnswerCountAtPostTime	Title
1	8704	2008-08-12	1042	08/11/2008 18:02:48	1	0	Final managed exception handler in a mixed na
2	8702	2008-08-12	1097	08/12/2008 11:40:09	1	0	What is the IDE for classic ASP and VBScript?
3	9009	2008-08-12	238	08/03/2008 21:42:37	94	13	Perl conditional operator is giving the wrong re
4	8728	2008-08-12	1008	08/11/2008 12:53:05	21	2	Delphi MDI Application and the titlebar of the N
5	8726	2008-08-12	572	08/06/2008 20:56:54	115	11	Checking FTP status codes with a PHP script.



Delta Lake 2.0 – All of Delta is now open

Unlock the power of Delta Lake

								
ACID Transactions	Scalable Metadata	Time Travel	Open Source	OPTIMIZE	OPTIMIZE ZORDER	Change data feed	Generated column support w/ partitioning	Coming Soon!
								
Unified Batch/Streaming	Schema Evolution /Enforcement	Audit History	DML Operations	Table Restore	S3 Multi-cluster writes	Data Skipping via Column Stats	Identity Columns	Iceberg to Delta converter
								
Compaction	MERGE Enhancements	Stream Enhancements	Simplified Logstore	Generated Columns	Multi-part checkpoint writes	Column Mapping	Subqueries in deletes and updates	Fast metadata only deletes





Delta Lake 2.0 – All of Delta is now open

Unlock the power of Delta Lake



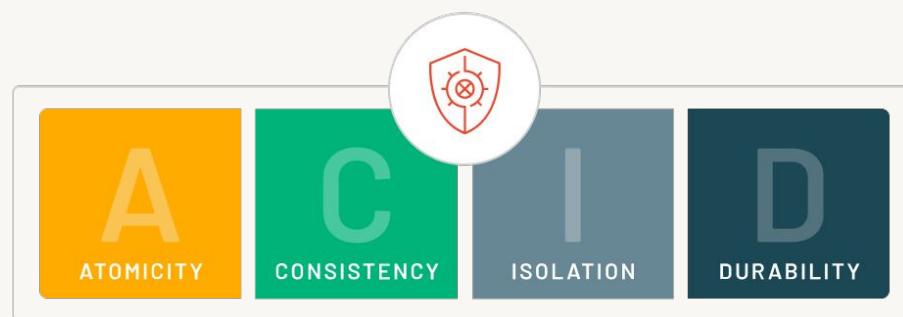
ACID
Transactions



Scalable
Metadata



OPTIMIZE
ZORDER



Data Skipping
via Column
Stats



Delta Lake File Structure

```
1 %fs ls /Users/vinny.vijeyakumaar@databricks.com/custom_demos/lakehouse/stackoverflow_train01/PostCreationDate=2008-07-31/
```

Python ► ▾ 🔍 ⌂ ⌂ ⌂

Table	+	Give feedback		
	path	name	size	modificationTime
1	dbfs:/Users/vinny.vijeyakumaar@databricks.com/custom_demos/lakehouse/stackoverflow_train01/PostCreationDate=2008-07-31/part-00000-28f21626-97da-4160-87e8-11cc19d91da3.c000.snappy.parquet	part-00000-28f21626-97da-4160-87e8-11cc19d91da3.c000.snappy.parquet	5614	1663966928000

Follow along @
<https://vnjv.co/dec22>



Delta Lake File Structure

```
1 %fs ls /Users/vinny.vijeyakumaar@databricks.com/custom_demos/lakehouse/stackoverflow_train01/_delta_log/
```

Table : +

	path	▲ name	▲ size	▲ modificationTime
2	dbfs:/Users/vinny.vijeyakumaar@databricks.com/custom_demos/lakehouse/stackoverflow_train01/_delta_log/.s3-optimization-1	.s3-optimization-1	0	1663889875000
3	dbfs:/Users/vinny.vijeyakumaar@databricks.com/custom_demos/lakehouse/stackoverflow_train01/_delta_log/.s3-optimization-2	.s3-optimization-2	0	1663889875000
4	dbfs:/Users/vinny.vijeyakumaar@databricks.com/custom_demos/lakehouse/stackoverflow_train01/_delta_log/00000000000000000000000000000000.crc	00000000000000000000000000000000.crc	3159	1663889875000
5	dbfs:/Users/vinny.vijeyakumaar@databricks.com/custom_demos/lakehouse/stackoverflow_train01/_delta_log/00000000000000000000000000000000.json	00000000000000000000000000000000.json	4129	1663889874000
6	dbfs:/Users/vinny.vijeyakumaar@databricks.com/custom_demos/lakehouse/stackoverflow_train01/_delta_log/00000000000000000000000000000001.crc	00000000000000000000000000000001.crc	3177	1663967267000
7	dbfs:/Users/vinny.vijeyakumaar@databricks.com/custom_demos/lakehouse/stackoverflow_train01/_delta_log/00000000000000000000000000000001.json	00000000000000000000000000000001.json	2445162	1663967264000

Let's generate more transactions



```
import random

for _ in range(250):
    [x, y] = random.sample(range(100000,999999), 2)
    spark.sql(f"""
        INSERT INTO vinny_vijeyakumar.stackoverflow_train01 VALUES
        ({x}, DATE("2007-08-01"), "", {x}, "", "", "", "", "", "", "", "", "", ""),
        ({y}, DATE("2007-08-01"), "", {y}, "", "", "", "", "", "", "", "", "", "")
    """)
```

A new log per successful transaction set

Displaying contents of /Users/vinny.vijeyakumaar@databricks.com/custom_demos/lakehouse/stackoverflow_train01/_delta_log/

Table : +

	path	name	size	modificationTime
515	dbfs:/Users/vinny.vijeyakumaar@databricks.com/custom_demos/lakehouse/stackoverflow_train01/_delta_log/0000000000000000254.json	00000000000000000000000000000000254.json	1869	1663983013000
516	dbfs:/Users/vinny.vijeyakumaar@databricks.com/custom_demos/lakehouse/stackoverflow_train01/_delta_log/00000000000000000000000000000000255.crc	00000000000000000000000000000000255.crc	3182	1663983017000
517	dbfs:/Users/vinny.vijeyakumaar@databricks.com/custom_demos/lakehouse/stackoverflow_train01/_delta_log/00000000000000000000000000000000255.json	00000000000000000000000000000000255.json	1869	1663983016000
518	dbfs:/Users/vinny.vijeyakumaar@databricks.com/custom_demos/lakehouse/stackoverflow_train01/_delta_log/00000000000000000000000000000000256.crc	00000000000000000000000000000000256.crc	3182	1663983019000
519	dbfs:/Users/vinny.vijeyakumaar@databricks.com/custom_demos/lakehouse/stackoverflow_train01/_delta_log/00000000000000000000000000000000256.json	00000000000000000000000000000000256.json	1869	1663983018000
520	dbfs:/Users/vinny.vijeyakumaar@databricks.com/custom_demos/lakehouse/stackoverflow_train01/_delta_log/_last_checkpoint	_last_checkpoint	6854	1663982904000

Peeking inside a log file

```
{  
  "protocol": {  
    "minReaderVersion": 1,  
    "minWriterVersion": 2  
  }  
}  
  
{  
  "metaData": {  
    "id": "00a4f716-7cf3-4c2a-8804-16bcb4ead707",  
    "format": {  
      "provider": "parquet",  
      "options": {}  
    },  
    "schemaString": "{\"type\":\"struct\", \"fields\": [ {\"name\":\"PostId\", \"type\":\"string\", \"nullable\":true, \"metadata\":{}}, {\"name\":\"PostCreationDate\", \"type\":\"date\", \"nullable\":true, \"metadata\":{}}, {\"name\":\"OwnerUserId\", \"type\":\"string\", \"nullable\":true, \"metadata\":{}}, {\"name\":\"OwnerCreationDate\", \"type\":\"string\", \"nullable\":true, \"metadata\":{}}, {\"name\":\"ReputationAtPostCreation\", \"type\":\"string\", \"nullable\":true, \"metadata\":{}}, {\"name\":\"OwnerUndeletedAnswerCountAtPostTime\", \"type\":\"string\", \"nullable\":true, \"metadata\":{}}, {\"name\":\"Title\", \"type\":\"string\", \"nullable\":true, \"metadata\":{}}, {\"name\":\"BodyMarkdown\", \"type\":\"string\", \"nullable\":true, \"metadata\":{}}, {\"name\":\"Tag1\", \"type\":\"string\", \"nullable\":true, \"metadata\":{}}, {\"name\":\"Tag2\", \"type\":\"string\", \"nullable\":true, \"metadata\":{}}, {\"name\":\"Tag3\", \"type\":\"string\", \"nullable\":true, \"metadata\":{}}, {\"name\":\"Tag4\", \"type\":\"string\", \"nullable\":true, \"metadata\":{}}, {\"name\":\"Tag5\", \"type\":\"string\", \"nullable\":true, \"metadata\":{}}, {\"name\":\"PostClosedDate\", \"type\":\"string\", \"nullable\":true, \"metadata\":{}}, {\"name\":\"OpenStatus\", \"type\":\"string\", \"nullable\":true, \"metadata\":{}} ] }",  
    "partitionColumns": [  
      "PostCreationDate"  
    ],  
    "configuration": {  
      "delta.autoOptimize.autoCompact": "true",  
      "delta.autoOptimize.optimizeWrite": "true"  
    },  
    "createdTime": 1663966838668  
  }  
}
```

Follow along @
<https://vnjv.co/dec22>



Log Contents

```
● ● ●  
{  
  "add": {  
    "path": "PostCreationDate=2008-08-03/part-00000-ddf05b95-9e25-4558-b3d7-  
96475df8362d.c000.snappy.parquet",  
    "partitionValues": {  
      "PostCreationDate": "2008-08-03"  
    },  
    "size": 22643,  
    "modificationTime": 1663966929000,  
    "dataChange": true,  
    "stats": "{\"numRecords\":39,\"minValues\":  
\\\"PostId\\\":\\\"588\\\",\\\"OwnerUserId\\\":\\\"101\\\",\\\"OwnerCreationDate\\\":\\\"08/01/2008  
12:01:23\\\",\\\"ReputationAtPostCreation\\\":\\\"1\\\",\\\"OwnerUndeletedAnswerCountAtPostTime\\\":\\\"0\\\",\\\"Title\\\":\\\"  
.NET Testing Framework Advice\\\",\\\"BodyMarkdown\\\":\\\"Alright, this is a rather odd  
qu\\\",\\\"Tag1\\\":\\\".net\\\",\\\"Tag2\\\":\\\".net\\\",\\\"Tag3\\\":\\\"api\\\",\\\"Tag4\\\":\\\"2008\\\",\\\"Tag5\\\":\\\"asp.net\\\",\\\"Post  
ClosedDate\\\":\\\"06/19/2012 17:52:00\\\",\\\"OpenStatus\\\":\\\"not constructive\\\"},\\\"maxValues\\\":  
\\\"PostId\\\":\\\"888\\\",\\\"OwnerUserId\\\":\\\"93\\\",\\\"OwnerCreationDate\\\":\\\"08/03/2008  
21:42:37\\\",\\\"ReputationAtPostCreation\\\":\\\"56\\\",\\\"OwnerUndeletedAnswerCountAtPostTime\\\":\\\"8\\\",\\\"Title\\\":\\\"  
cx_Oracle - what is the best way\\ufffd\\\",\\\"BodyMarkdown\\\":\\\"[Django][1] view points to a  
fun\\ufffd\\\",\\\"Tag1\\\":\\\"windows\\\",\\\"Tag2\\\":\\\"windows\\\",\\\"Tag3\\\":\\\"xslt\\\",\\\"Tag4\\\":\\\"vb.net\\\",\\\"Tag5\\\":\\\"  
zend\\\",\\\"PostClosedDate\\\":\\\"06/19/2012 17:52:00\\\",\\\"OpenStatus\\\":\\\"open\\\"},\\\"nullCount\\\":  
\\\"PostId\\\":0,\\\"OwnerUserId\\\":0,\\\"OwnerCreationDate\\\":0,\\\"ReputationAtPostCreation\\\":0,\\\"OwnerUndeleted  
AnswerCountAtPostTime\\\":0,\\\"Title\\\":0,\\\"BodyMarkdown\\\":0,\\\"Tag1\\\":0,\\\"Tag2\\\":3,\\\"Tag3\\\":12,\\\"Tag4\\\":26,  
\\\"Tag5\\\":34,\\\"PostClosedDate\\\":38,\\\"OpenStatus\\\":0}}},  
    "tags": {  
      "INSERTION_TIME": "1663966928000003",  
      "MIN_INSERTION_TIME": "1663966928000003",  
      "MAX_INSERTION_TIME": "1663966928000003",  
      "OPTIMIZE_TARGET_SIZE": "268435456"  
    }  
  }  
}
```

We know this is a reliable, fully-written out file, as the outcome of a successful DML transaction

Column stats are useful during query planning time to determine which files can be skipped (skipped files → less I/O, less memory required, reduced chance of disk spill)

Log entry after a MERGE INTO ... command

```
● ● ●

{
  "remove": {
    "path": "PostCreationDate=2007-08-01/part-00000-f982e89f-4264-43bb-a043-
f1ba0a27492c.c000.snappy.parquet",
    "deletionTimestamp": 1664314049556,
    "dataChange": true,
    "extendedFileMetadata": true,
    "partitionValues": {
      "PostCreationDate": "2007-08-01"
    },
    "size": 3376,
    "tags": {
      "INSERTION_TIME": "1664314022000000",
      "MIN_INSERTION_TIME": "1664314022000000",
      "MAX_INSERTION_TIME": "1664314022000000",
      "OPTIMIZE_TARGET_SIZE": "268435456"
    }
  }
}

{
  "add": {
    "path": "PostCreationDate=2007-08-01/part-00000-5d440efc-3877-4fa4-8486-
6725dc987151.c000.snappy.parquet",
    "partitionValues": {
      "PostCreationDate": "2007-08-01"
    },
    "size": 3550,
    "modificationTime": 1664314050000,
    "dataChange": true,
    "stats": "{\"numRecords\":4,\"minValues\":
{\\"PostId\\":\"1078360\",\\\"OwnerUserId\\\":\\\"\",\\\"OwnerCreationDate\\\":\\\"1078360\\\",\\\"ReputationAtPostCreation\\\":\\\"\",\\\"OwnerUndeletedAnswerCountAtPostTime\\\":\\\"\",\\\"Title\\\":\\\"\",\\\"BodyMarkdown\\\":\\\"\",\\\"Tag1\\\":\\\"\",\\
\"Tag2\\\":\\\"\",\\\"Tag3\\\":\\\"\",\\\"Tag4\\\":\\\"\",\\\"Tag5\\\":\\\"\",\\\"PostClosedDate\\\":\\\"\",\\\"OpenStatus\\\":\\\"\"},\\\"max
Values\\\":
{\\"PostId\\\":\\\"999999\\\",\\\"OwnerUserId\\\":\\\"\",\\\"OwnerCreationDate\\\":\\\"999999\\\",\\\"ReputationAtPostCreation\\
:\\\"\",\\\"OwnerUndeletedAnswerCountAtPostTime\\\":\\\"\",\\\"Title\\\":\\\"\",\\\"BodyMarkdown\\\":\\\"\",\\\"Tag1\\\":\\\"\",\\\"T
```



Unpacking "stats"

```
{  
    "numRecords": 5,  
    "minValues": {  
        "PostId": "11",  
        "OwnerUserId": "1",  
        "OwnerCreationDate": "07/31/2008 14:22:31",  
        "ReputationAtPostCreation": "1",  
        "OwnerUndeletedAnswerCountAtPostTime": "0",  
        "Title": "Decimal vs Double?",  
        "BodyMarkdown": "Are there any conversion tools f",  
        "Tag1": "c#",  
        "Tag2": "css",  
        "OpenStatus": "open"  
    },  
    "maxValues": {  
        "PostId": "9",  
        "OwnerUserId": "9",  
        "OwnerCreationDate": "07/31/2008 21:35:26",  
        "ReputationAtPostCreation": "16",  
        "OwnerUndeletedAnswerCountAtPostTime": "2",  
        "Title": "Tools for porting J# code to C#",  
        "BodyMarkdown": "I've got an absolutely positione\ufffd",  
        "Tag1": "j#",  
        "Tag2": "css",  
        "OpenStatus": "open"  
    },  
    "nullCount": {  
        "PostId": 0,  
        "OwnerUserId": 0,  
        "OwnerCreationDate": 0,  
        "ReputationAtPostCreation": 0,  
        "OwnerUndeletedAnswerCountAtPostTime": 0,  
        "Title": 0,  
        "BodyMarkdown": 0,  
        "Tag1": 0,  
        "Tag2": 4,  
        "Tag3": 5,  
        "Tag4": 5,  
        "Tag5": 5,  
        "PostClosedDate": 5,  
        "OpenStatus": 0  
    }  
}
```

```
{  
    "numRecords": 5,  
    "minValues": {  
        "PostId": "11",  
        "OwnerUserId": "1",  
        "OwnerCr  
        "Reputat  
        "OwnerUn  
        "Title":  
        "BodyMar  
        "Tag1":  
        "Tag2":  
        "OpenSta  
    },  
    "maxValues": {  
        "PostId": "9",  
        "OwnerUserId": "9",  
        "OwnerCreationDate": "07/31/2008 21:35:26",  
        "ReputationAtPostCreation": "16",  
        "OwnerUndeletedAnswerCountAtPostTime": "2",  
        "Title": "Tools for porting J# code to C#",  
        "BodyMarkdown": "I've got an absolutely positione\ufffd",  
        "Tag1": "j#",  
        "Tag2": "css",  
        "OpenStatus": "open"  
    }  
}
```

```
{  
    "maxValues": {  
        "PostId": "9",  
        "OwnerUserId": "9",  
        "OwnerCreationDate": "07/31/2008 21:35:26",  
        "ReputationAtPostCreation": "16",  
        "OwnerUndeletedAnswerCountAtPostTime": "2",  
        "Title": "Tools for porting J# code to C#",  
        "BodyMarkdown": "I've got an absolutely positione\ufffd",  
        "Tag1": "j#",  
        "Tag2": "css",  
        "OpenStatus": "open"  
    }  
}
```

```
{  
    "nullCount": {  
        "PostId": 0,  
        "OwnerUserId": 0,  
        "OwnerCreationDate": 0,  
        "ReputationAtPostCreation": 0,  
        "OwnerUndeletedAnswerCountAtPostTime": 0,  
        "Title": 0,  
        "BodyMarkdown": 0,  
        "Tag1": 0,  
        "Tag2": 4,  
        "Tag3": 5,  
        "Tag4": 5,  
        "Tag5": 5,  
        "PostClosedDate": 5,  
        "OpenStatus": 0  
    }  
}
```

ZORDER BY (col1, ...)

Read in each part-file and test for $X = 2$ or $Y = 3$

Linear Order

9 files scanned in total
21 false positives



Z-order

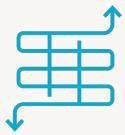
7 files scanned in total
13 false positives





Delta Lake 2.0 – All of Delta is now open

Unlock the power of Delta Lake



Schema Evolution
/Enforcement



OOTB Schema Enforcement

```
1 %sql
2 INSERT INTO vinny_vijeyakumaar.stackoverflow_train01 VALUES
3   ("abc", "2007-08-01", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "extraColumn"),
4   ("xyz", DATE("2007-08-01"), "", 999999, "", "", "", "", "", "", "", "", "", "", "", "", "")
```

⊕**AnalysisException:** expected 17 columns but found 16 columns in row 1; line 1 pos 53

OOTB Schema Enforcement

```
1 %sql
2 INSERT INTO vinny_vijeyakumaar.stackoverflow_train01 VALUES
3   ("abc", "2007-08-01", "", "", "", "", "", "", "", "", "", "", "", "", ""),
4   ("xyz", DATE("2007-08-01"), "", 999999, "", "", "", "", "", "", "", "", "", "", "")
```

⊕**AnalysisException:** incompatible types found in column col2 for inline table; line 1 pos 53

Adding Constraints



```
ALTER TABLE vinny_vijeyakumar.stackoverflow_train01  
ADD CONSTRAINT post_id_not_null CHECK (PostId IS NOT NULL);
```

```
ALTER TABLE vinny_vijeyakumar.stackoverflow_train01  
ADD CONSTRAINT post_creation_date_valid CHECK (  
PostCreationDate BETWEEN "2007-08-01" AND "2012-09-01");
```

Constraint Enforcement

```
1 %sql
2 INSERT INTO vinny_vijeyakumaar.stackoverflow_train01 VALUES
3 ("xyz", DATE("2022-09-29"), "", 99999, "", "", "", "", "", "", "", "", "", "", "")
```

▶ (1) Spark Jobs

```
✉ com.databricks.sql.transaction.tahoe.schema.DeltaInvariantViolationException: CHECK constraint post_creation
_date_valid ((PostCreationDate >= '2007-08-01') AND (PostCreationDate <= '2012-09-01')) violated by row with
values:
```

Follow along @
<https://vnjv.co/dec22>





Delta Lake 2.0 – All of Delta is now open

Unlock the power of Delta Lake



Time Travel



Audit History



Table Restore



**WHERE WE'RE GOING WE
DON'T NEED...ROADS**





```
DESCRIBE HISTORY vinny_vijeyakumaar.stackoverflow_train01  
LIMIT 10 -- get the last 10 operations only
```

Table							+	Give feedback
	version	timestamp	userId	userName	operation	operationParameters		
1	274	2022-09-28T02:03:37.000+0000	5211215277230706	vinny.vijeyakumar@databricks.com	MERGE	▶ {"predicate": "(target.PostId = source \"notMatchedPredicates": [{"actionType": "MERGE", "key": "PostId", "value": 123}], "numTargetRowsCopied": 196, "numTargetRowsDeleted": 0, "numTargetRowsInserted": 5529, "numTargetRowsUpdated": 200, "numTargetChangeFilesAdded": 1, "numSourceRows": 4, "numTargetFiles": 1087}"}		
2	273	2022-09-27T23:25:04.000+0000	5211215277230706	vinny.vijeyakumar@databricks.com	WRITE	▶ {"mode": "Append", "partitionBy": []}		
3	272	202	clusterId	readVersion	isolationLevel	isBlindAppend	operationMetrics	
4	271	202	0923-094955-git4n0b4	273	WriteSerializable	false	▶ {"numTargetRowsCopied": 196, "numTargetRowsDeleted": 0, "numTargetRowsInserted": 5529, "numTargetRowsUpdated": 200, "numTargetChangeFilesAdded": 1, "numSourceRows": 4, "numTargetFiles": 1087}	
5	270	202						
6	269	202						
7	268	202	0923-094955-git4n0b4	272	WriteSerializable	true	▶ {"numFiles": 1, "numOutputRows": 2, "numOutputBytes": 3228}	
8	267	202	0923-094955-git4n0b4	271	WriteSerializable	true	▶ {"numFiles": 1, "numOutputRows": 2, "numOutputBytes": 3228}	
9	266	202	0923-094955-git4n0b4	270	WriteSerializable	true	▶ {"numFiles": 1, "numOutputRows": 2, "numOutputBytes": 3226}	
			0923-094955-git4n0b4	269	WriteSerializable	true	▶ {"numFiles": 1, "numOutputRows": 2, "numOutputBytes": 3228}	
			0923-094955-git4n0b4	268	WriteSerializable	true	▶ {"numFiles": 1, "numOutputRows": 2, "numOutputBytes": 3228}	
			0923-094955-git4n0b4	267	WriteSerializable	false	▶ {"numRemovedFiles": 1, "numCopiedRows": 85, "numAddedChangeFiles": 1, "scanTimeMs": 1610, "numAddedFiles": 1, "numUpdatedRows": 2, "rewriteType": "MERGE"}	
			0923-094955-git4n0b4	266	WriteSerializable	false	▶ {"numRemovedFiles": 1, "numCopiedRows": 93, "numAddedChangeFiles": 1, "scanTimeMs": 1784, "numAddedFiles": 1, "numUpdatedRows": 2, "rewriteType": "MERGE"}	
©2022 Databricks	0923-094955-git4n0b4	265		WriteSerializable	true		▶ {"numFiles": 1, "numOutputRows": 2, "numOutputBytes": 3228}	

Time Travel by version number or timestamp



```
WITH rec_count_00 AS (
    SELECT COUNT(*) AS n FROM vinny_vijeyakumar.stackoverflow_train01
    VERSION AS OF 1
)
, rec_count_01 AS (
    SELECT COUNT(*) AS n FROM vinny_vijeyakumar.stackoverflow_train01
    TIMESTAMP AS OF (CURRENT_TIMESTAMP() - INTERVAL 240 MINUTES)
)
, rec_count_now AS (
    SELECT COUNT(*) AS n FROM vinny_vijeyakumar.stackoverflow_train01
)
SELECT r0.n AS start, r1.n AS 240minsAgo, rn.n AS now
FROM rec_count_00 r0, rec_count_01 r1, rec_count_now rn
```

Follow along @
<https://vnjv.co/dec22>



start	240minsAgo	now
3352998	3353020	3353280

Restore to a previous state

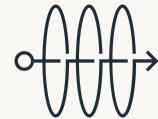


```
-- Restore table to previous state
RESTORE TABLE vinny_vijeyakumaar.stackoverflow_train01
TO VERSION AS OF 12
-- TO TIMESTAMP AS OF "2022-09-28 10:00:00"
```



Delta Lake 2.0 – All of Delta is now open

Unlock the power of Delta Lake



Change data
feed



Audit History



Track changes to your data with Change Data Feed (CDF)

```
1 %sql
2 SELECT * FROM TABLE_CHANGES("vinny_vijeyakumaar.stackoverflow_train01", 270)
3 WHERE PostId = "13211"
```

▶ (3) Spark Jobs

Table : +							Give feedback
	Tag5	PostClosedDate	OpenStatus	_change_type	_commit_version	_commit_timestamp	
1	null	11/21/2011 07:45:23	not constructive	update_postimage	289	2022-09-25T01:24:59.000+0000	
2	systems	11/21/2011 07:45:23	not constructive	update_preimage	289	2022-09-25T01:24:59.000+0000	

Propagate changes downstream



```
INSERT INTO target_schema.target_table
SELECT * EXCEPT (_change_type, _commit_version, _commit_timestamp)
-- Let's assume our orchestration tool knows that the last update happened at 2022-09-29 00:00:00
FROM TABLE_CHANGES("vinny_vijeyakumaar.stackoverflow_train01", "2022-09-29 00:00:00")
WHERE _change_type IN ("insert", "update_postimage");

DELETE FROM target_schema.target_table
WHERE PostId IN (
    SELECT PostId
    -- Let's assume our orchestration tool knows that the last update happened at 2022-09-29 00:00:00
    FROM TABLE_CHANGES("vinny_vijeyakumaar.stackoverflow_train01", "2022-09-29 00:00:00")
    WHERE _change_type IN ("delete")
);
```



Delta Lake 2.0 – All of Delta is now open

Unlock the power of Delta Lake



ACID
Transactions



Scalable
Metadata



Time Travel



Open Source



OPTIMIZE



Change data
feed



Schema Evolution
/Enforcement



Audit History



DML Operations



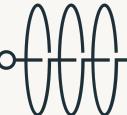
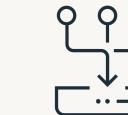
Data Skipping
via Column
Stats





Delta Lake 2.0 – All of Delta is now open

Unlock the power of Delta Lake

								
ACID Transactions	Scalable Metadata	Time Travel	Open Source	OPTIMIZE	OPTIMIZE ZORDER	Change data feed	Generated column support w/ partitioning	Coming Soon!
								
Unified Batch/Streaming	Schema Evolution /Enforcement	Audit History	DML Operations	Table Restore	S3 Multi-cluster writes	Data Skipping via Column Stats	Identity Columns	Iceberg to Delta converter
								
Compaction	MERGE Enhancements	Stream Enhancements	Simplified Logstore	Generated Columns	Multi-part checkpoint writes	Column Mapping	Subqueries in deletes and updates	Fast metadata only deletes



Deep-dive Resources

Follow along @
<https://vnjv.co/dec22>



**Diving Into Delta Lake:
Unpacking The Transaction
Log**

by Burak Yavuz, Michael Armbrust and Brenner Heintz
August 21, 2019 in Company Blog

Transaction Log
Single Commits
(Optional) Partition Directories
Data Files

```
my_table/
  _delta_log/
    0000.json
    0001.json
  date=2019-01-01/
    file-1.parquet
```

[Link](#)

**Diving Into Delta Lake:
Schema Enforcement &
Evolution**

by Burak Yavuz and Brenner Heintz
September 24, 2019 in Company Blog

```
# Add the mergeSchema option
loans.write.format("delta") \
  .option("mergeSchema", "true") \
  .mode("append") \
  .save(DELTALAKE_SILVER_PATH)
```

[Link](#)

**Diving Into Delta Lake: DML
Internals (Update, Delete,
Merge)**

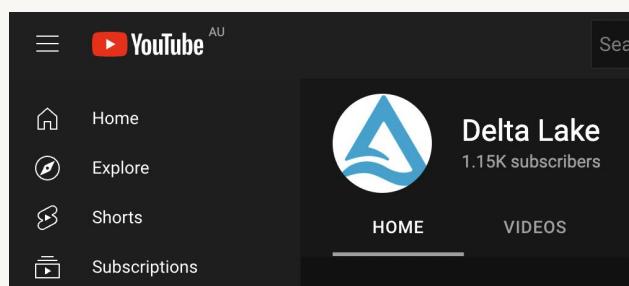
by Tathagata Das and Brenner Heintz
September 29, 2020 in Engineering Blog

Merge – Under the hood

Scan 1: Inner join between target and source to select files that have matches

Scan 2: Outer join between the selected files in target and source and write the update/deleted/inserted data

[Link](#)



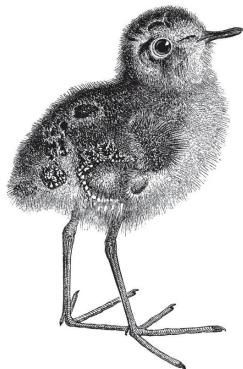
[YouTube Channel](#): Deep Dives, AMAs, Community Office Hours, Tutorials, Tech Talks

O'REILLY®

Delta Lake The Definitive Guide

Modern Data Lakehouse Architectures
with Delta Lake

Early
Release
Raw & Unedited
Sponsored by
databricks



Denny Lee,
Tathagata Das &
Vini Jaiswal

[Free preview!](#)

The cover features the title 'Simplifying Data Engineering and Analytics with Delta' in large white font. Below it is a subtitle 'Create analytics-ready data that fuels artificial intelligence and business intelligence'. The background is dark with abstract data visualization elements like numbers and lines.

Anindita Mahapatra
Foreword by Doug May, VP, Global Value Acceleration - Databricks Inc.

[Amazon](#)

The cover has a dark background with the title 'Delta Lake: High-Performance ACID Table Storage over Cloud Object Stores' in white. Below it is a short abstract and author information. The bottom half contains dense text from the paper.

Delta Lake: High-Performance ACID Table Storage over Cloud Object Stores

Michael Armbrust, Tathagata Das, Liwen Sun, Burak Yavuz, Shixiong Zhu, Mukul Murthy, Joseph Torres, Herman van Hovell, Adrian Ionescu, Alicja Luszczak, Michał Świątkowski, Michał Szafrański, Xiao Li, Takuya Ueshim, Mostafa Mokhtar, Peter Boncz¹, Ali Ghodsi¹, Sameer Paranjape¹, Prashant Raghav¹, Ronald Xin¹, Matei Zaharia¹
Databricks¹, CWI¹, UC Berkeley¹, Stanford University¹
delta-paper-authors@databricks.com

ABSTRACT
Cloud object stores such as Amazon S3 are some of the largest and most cost-effective storage systems on the planet, making them an attractive target to store large data warehouses and data lakes. Unfortunately, they are known to have serious shortcomings that make it difficult to achieve ACID transactions and high performance metadata operations such as listing objects are expensive, and consistency guarantees are limited. In this paper, we present Delta Lake, an open source ACID table store that builds upon object stores originally developed at Databricks. Delta Lake uses a transaction log that is compacted into Apache Parquet format to provide ACID properties, time travel, and significant updates to provide ACID properties, time travel, and significant updates to provide ACID properties, table-level atomicity, and the ability to quickly search billions of table partitions for those relevant to a query. It also leverages this design to support high-level features such as automatic data layout optimization, upserts, and materialized views. Delta Lake tables can be accessed from Apache Spark, Hive, Presto, Redshift and other systems. Delta Lake is deployed at thousands of Databricks customers that process exabytes of data per day, with the largest instances processing petabytes of data and billions of objects.

VLDB Reference Format:
Armbrust et al. Delta Lake: High-Performance ACID Table Storage over Cloud Object Stores. *VLDB*, 13(12): 3411–3424, 2020.
DOI: <https://doi.org/10.14778/3415478.3415560>

1. INTRODUCTION
Cloud object stores such as Amazon S3 [3] and Azure Blob Storage [17] have become one of the largest and most widely used storage systems on the planet, holding exabytes of data for millions of customers [46]. Apart from the traditional advantages of clouds services, such as pay-as-you-go billing, economies of scale, and expert management [51], cloud object stores are also attractive because they allow for scale computation and storage resources separately: for example, a user can store a petabyte of data but only run a cluster to execute a query over it for a few hours.

In our experience working with cloud customers, these consistency and performance issues create major challenges for enterprise data teams. Most enterprise datasets are inherently updateable, so they require a solution for updates with most data being about users: require table-wide updates to implement privacy policies such as GDPR compliance [27]; and even purely internal datasets may require updates to repair inconsistent data, incorporate late records, etc. As a result, in the last few years (from 2014 to 2016), around half the support escalations we received were due to data corruption, consistency or performance issues due to cloud storage strategies (e.g., undergoing the cost of a crashed update job, or impacting the performance of a query that reads tens of thousands of objects).

[VLDB Paper](#)

Join the community today!



delta.io



go.delta.io/github



go.delta.io/slack



go.delta.io/twitter

Delta Lake Roadmap

Delta Lake [project's roadmap](#) is publicly available [in this GitHub issue](#)

Key upcoming features



Support for
Spark 3.3



Extended Flink
Support



Extended Rust
support

VISIT: <https://delta.io/roadmap/>

Join us!



DATA+AI
WORLD TOUR 2022

Destination Lakehouse

IN PERSON + VIRTUAL

SYDNEY | 18 NOV
VIRTUAL | 1 DEC



Free to attend

ORGANIZED BY  databricks

Coming The Star Sydney 17 & 18 Nov



Matei Zaharia

Co-founder and Chief
Technologist, Databricks;
Original Creator of Apache
Spark™ and MLflow

DATABRICKS

- **SKILL-UP** with complimentary training and onsite certification
- **EXPLORE** the technical and business value of the lakehouse with industry experts featuring Databricks Co-founder, Matei Zaharia
- **DEEP DIVE** into the latest innovations on the Databricks Lakehouse Platform through practical use cases and real customer scenarios shared by **Atlassian, Alinta Energy, AEMO, Sportsbet**, and many others



Free to attend



Delta Lake: Turn your Data Lake into a Lakehouse

Vinoaj (Vinny) Vijeyakumaar
Senior Solutions Architect
vinny@databricks.com |

©2022 Databricks Inc. — All rights reserved

/vinoaj

